



## ESiWACE3 - Excellence in Simulation of Weather and Climate in Europe, Phase 3

# GPU Optimization with Kernel Tuner

September 13, 2024



**EuroHPC**  
Joint Undertaking

*Funded by the European Union. This work has received funding from the European High Performance Computing Joint Undertaking (JU) under grant agreement No 101093054.*



Co-funded by  
the European Union

netherlands

**eScience** center

# Outline for the day

---

10:00	–	10:15	Welcome
10:15	–	10:45	Energy efficient GPU computing
10:45	–	11:30	Code optimizations for energy efficiency
11:30	–	11:45	Break
11:45	–	12:15	Mixed-precision programming techniques
12:15	–	12:45	Optimizing GPU core clock frequency
12:45	–	13:00	Q&A

---

# Administrative announcements

- We will have four sessions in which we start with introducing some new concepts and follow with a hands-on exercise
- The hands-on exercises include example kernels, but you are also welcome to experiment with your own code
- We will use JupyterHub hosted by the VSC for the hands-on, so you don't need to have a GPU or install anything locally
- You can download the slides here:
  - [https://github.com/KernelTuner/kernel\\_tuner\\_tutorial/tree/master/slides/2024\\_VSC\\_ESiWACE3](https://github.com/KernelTuner/kernel_tuner_tutorial/tree/master/slides/2024_VSC_ESiWACE3)

# Learning objectives

- Understand the energy footprint of computing
- Optimize applications for performance to reduce energy consumption
- Reduce data movement with mixed-precision techniques
- Tune GPU core frequencies to find the most energy-efficient configuration

# **Energy Efficient GPU Computing**

# LLM Training emissions

## CO2 equivalent emissions (tonnes) by select machine learning models and real-life examples, 2020–23

Source: AI Index, 2024; Luccioni et al., 2022; Strubell et al., 2019 | Chart: 2024 AI Index report




Figure 2.13.1

## What is 500 tons of CO<sub>2</sub>?

Roughly equal to:

- 8,268 tree seedlings grown for 10 years
- \$80,000 in electricity bill
- 63 homes' energy use for a year in the US
- 111 passenger cars driving around for a year in the US
  
- Less than 2 days of running the Frontier supercomputer ...



# Energy cost of supercomputers

Frontier: #1 in TOP500 list (Jun 2024)

- #13 Green500 (Jun 2024)
- 20 Megawatt continuously
- \$40 million annual electricity bill
- 100,000 metric tons of CO<sub>2</sub> annually
- ~20,000 cars on the road for a year in US

Summit: (#9, Frontier's predecessor)

- 64% of energy is consumed by GPUs




*Efficient Computation through Tuned Approximation*  
David Keyes, SIAG/SC Supercomputing Spotlights 2022


*Autotuning based on frequency scaling toward energy efficiency of blockchain algorithms on graphics processing units* M. Stachowski, A. Fiebig, and T. Rauber, Journal of Supercomputing, 2020.



## GFLOPs/W for different architectures




# Energy Efficiency of GPUs



# Energy, Heat, and Surface Size

- **Nvidia H100 GPU:**

- Energy: 350 Watt
- Surface: 8.14 cm<sup>2</sup>
- Heat dissipation: 43.0 Watt/cm<sup>2</sup>



- **Light bulb:**

- Energy: 100 Watt
- Surface: 15 cm<sup>2</sup>
- Heat dissipation: 6.7 Watt/cm<sup>2</sup>




- **Electric cooker:**

- Energy: 1800 Watt
- Surface: 1017 cm<sup>2</sup>
- Heat dissipation: 1.8 Watt/cm<sup>2</sup>



# GPU Temperature – Energy Efficiency relation



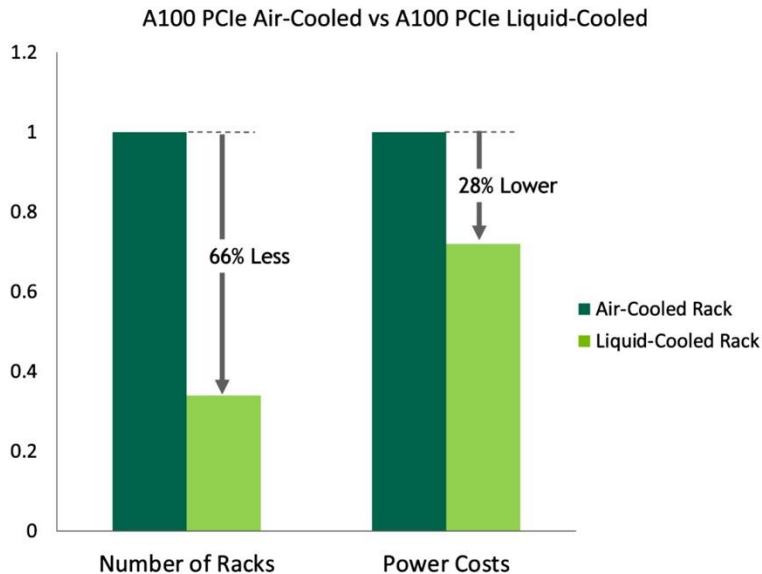
Hotter GPUs can be ~7%  
less energy efficient

Results obtained with xGPU  
(radio astronomy correlator)  
on Nvidia K20

## What about cooling?


- Liquid cooling is more energy friendly than air cooling
- But as the efficiency difference between hot and cold GPUs is ~7%, you probably shouldn't overdo the cooling

## RACK LEVEL COST REDUCTION



Configuration:  
2000 servers each with 2x CPU | 192GB | 1TB SSD | 2x A100 80GB  
Air-cooled and liquid-cooled GPUs each at 300W TDP and same performance characteristics  
Air-cooled infrastructure @ 1.6 PUE; Liquid-cooled infrastructure @ 1.15 PUE  
15KW Air-Cooled Rack | 30KW Liquid-Cooled Rack | Power costs = \$0.2 per KWhr

# Clock frequency power relation




# How is energy spent within a GPU?

Moving data around is 20x more expensive  
than computing on it

Estimations for Nvidia H100:

- A single double-precision Fused Multiply-Add<sup>1</sup>: 13.7 pJ
- Moving the operands (4x 64-bits) for 10 mm within chip<sup>2</sup>: 294.4 pJ (21x more energy)

```
mad.f64 %f1, %f2, %f3, %f0; // c += a*b;
```



<sup>1</sup>Based on 25.6 TFLOP/s peak at 350 Watt TDP <sup>2</sup>Based on 115 femtojoule/bit/mm  
<https://www.nvidia.com/en-us/data-center/h100/> *GPUs and the Future of Parallel Computing*  
Keckler et al. 2011

# How do we create energy efficient GPU applications?

Three strategies for energy efficient GPU Computing:

1. Use for shorter amount of time
2. Minimize data movements
3. Optimize device settings



# How do we create energy efficient GPU applications?

Three strategies for energy efficient GPU Computing:

1. Use for shorter amount of time → Optimize application performance
2. Minimize data movements → Lower/mixed precision techniques
3. Optimize device settings → Optimize clock frequency

# **Code Optimizations for Energy Efficiency**

# GPU code optimizations

- Modify the kernel source code to improve **performance**
- Effects on performance can be different on different GPUs or different input data
- You can tune:
  - Enabling or disabling an optimization
  - The parameters introduced by certain optimizations
- You often need to combine multiple different optimizations with specific tunable parameter values to arrive at optimal performance

## Further reading

- In March 2023, we published a literature review summarizing the last decade of code optimizations for GPU programming
  - We describe which optimizations are used in literature and how they are used
- *Optimization Techniques for GPU Programming*  
Pieter Hijma, Stijn Heldens, Alessio Sclocco, Ben van Werkhoven, and Henri Bal  
ACM Computing surveys 2023  
<https://dl.acm.org/doi/abs/10.1145/3570638>

# Overview of GPU Optimizations

- Coalescing memory accesses
- Host/device communication
- Kernel fusion
- Loop blocking
- Loop unrolling
- Prefetching
- Recomputing values
- Reducing atomics
- Reducing branch divergence
- Reducing redundant work
- Reducing register usage
- Reformatting input data
- Using a specific memory space
- Using warp shuffle instructions
- Varying work per thread
- Vectorization

# Overview of GPU Optimizations

- Coalescing memory accesses
- Host/device communication
- **Kernel fusion**
- **Loop blocking**
- Loop unrolling
- Prefetching
- Recomputing values
- Reducing atomics
- Reducing branch divergence
- Reducing redundant work
- Reducing register usage
- Reformatting input data
- Using a specific memory space
- Using warp shuffle instructions
- Varying work per thread
- Vectorization

# Kernel fusion

Merge one or more kernels into one kernel

- Why?
  - Reduces data movements between off-chip DRAM and GPU registers
  - Moving data around is more expensive than computing on it
- How?
  - Fuse the kernel arguments and computations of two kernels into one
  - Demote a kernel to a `__device__` function and call it from another kernel
  - *Temporal fusion*: merge multiple calls of the same kernel into one

# Kernel fusion

```
// c = a+b  
vector_add<<<grid, threads>>>(c, a, b, n);  
// e = c+d  
vector_add<<<grid, threads>>>(e, c, d, n);
```

```
__global__  
void vector_add(float *c, float *a, float *b,  
               int n) {  
    int i = (blockIdx.x*blockDim.x)+threadIdx.x;  
    if (i < n) {  
        c[i] = a[i] + b[i];  
    }  
}
```



```
// e = a+b+d  
vector_3add<<<grid, threads>>>(e, a, b, d, n);
```

```
__global__  
void vector_3add(float *d, float *a, float *b,  
                float *c, int n) {  
    int i = (blockIdx.x*blockDim.x)+threadIdx.x;  
    if (i < n) {  
        d[i] = a[i] + b[i] + c[i];  
    }  
}
```




# Loop blocking


Modify the structure of one or more loops to work in blocks over the data

- Why?
  - Increases spatial / temporal locality
  - Reduces the 'working set' of the algorithm
- How?
  - Change the order of computations and data accesses in nested loops
  - Usually nearly doubles the number of for-loops in the code
  - Outer-loops iterate over the blocks
  - Inner-loops iterate within each block

## Loop blocking



```
for (int j=0; j<ny; j++) {
    for (int i=0; i<nx; i++) {
        ...[j*nx + i]
    }
}
```



```
for (int j=0; j<ny; j+=nyb) {
    for (int i=0; i<nx; i+=nxb) {
```

```
        for (int jb=0; jb<nyb; jb++) {
            for (int ib=0; ib<nxb; ib++) {
                ...[(j+jb)*nx + (i+ib)]
            }
        }
    }
}
```

# Hands-on

# First hands-on

- The first hands-on notebook is:
  - [https://github.com/KernelTuner/kernel\\_tuner\\_tutorial/blob/master/hands-on/esiwace3/04\\_Code\\_Optimizations\\_for\\_Energy.ipynb](https://github.com/KernelTuner/kernel_tuner_tutorial/blob/master/hands-on/esiwace3/04_Code_Optimizations_for_Energy.ipynb)
- The goal of this hands-on is to:
  - See an example of **Kernel Fusion**
  - Compare the **energy consumption** of different kernels
- Please follow the instructions in the Notebook
- Feel free to ask questions to instructors and mentors

# **Mixed-Precision Programming Techniques**

# Low precision computing

**Double (64 bit)**



**Single (32 bit)**



**FP24 (24 bit)**



**TensorFloat (19 bit)**



**Half (16 bit)**



**Bfloat (16 bit)**



**Minifloat (8 bit)**



# Low precision computing

- Low precision has many benefits 😊!
  - Faster **computation**
    - Less compute cycles required, especially double precision is often slow
  - Lower **memory footprint**
    - Less bits required per number
  - Better **cache utilization**
    - Higher cache hit rates
  - Higher *effective* **memory bandwidth**
    - More numbers per second
  - Lower **register usage**
    - Increases GPU occupancy, thus performance
  - All these points also increase **energy efficiency**

# Benefits of low precision computing

- But, at the cost of loss in precision 😞

Floating-point	Value of Pi	Error
<i>Infinite bits</i>	3.141592653589793238462643383279502884197169399375...	0
<b>64 bit</b>	3.141592653589793115997963468544185161590576171875	$3.9 \times 10^{-15} \%$
<b>32 bit</b>	3.1415925025939941406250000000000000000000000000000	0.000005%
<b>16 bit</b>	3.140625000	0.03%
<b>8 bit</b>	3.000	4.5%



# Mixed-precision arithmetic

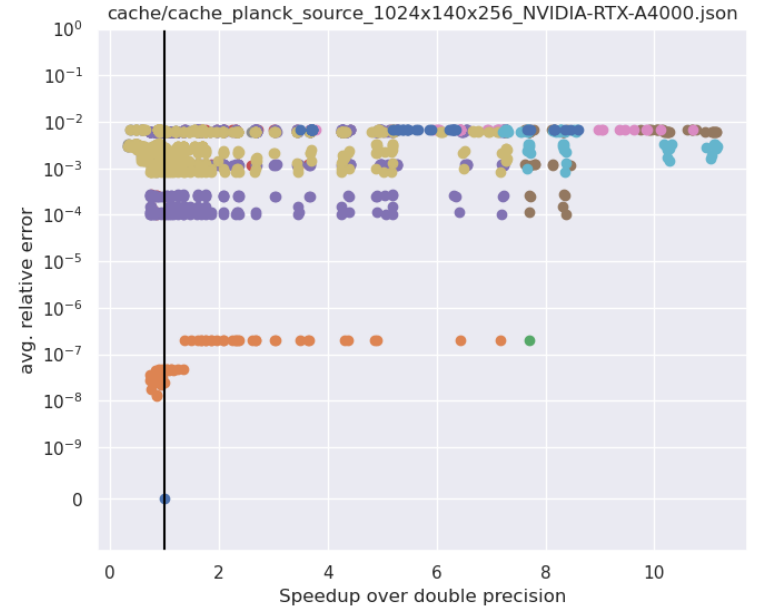
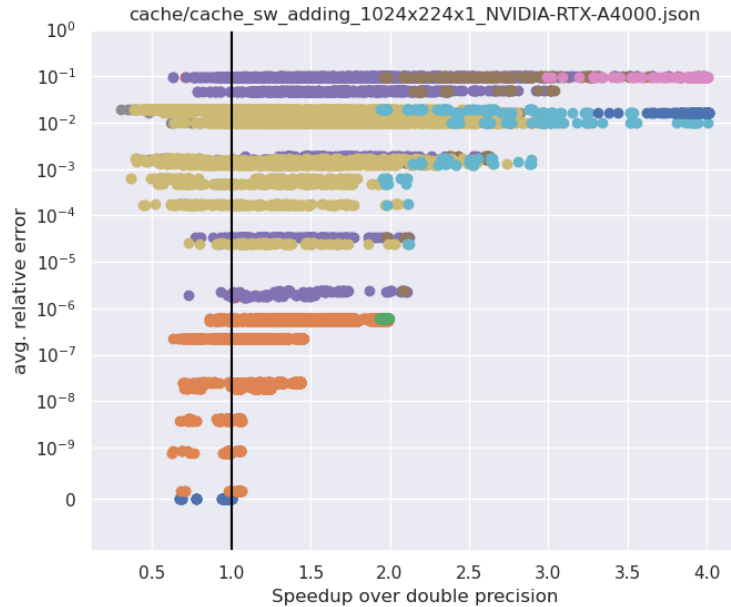
- Core idea of **mixed precision**:
  - What if we **mix different** precision levels in one application?
  - Use **different floating-point types** for **different variables** in code
- Trade-off between **performance** and **numerical accuracy**
  - **Lower precision** typically results in **higher performance**
  - Need to find **balance** between **error** and **speedup**
- Creates **huge** search space
  - What precision should be used for each variable?
  - Example: 20 variables and 4 precisions gives  $4^{20} = 1$  trillions combinations!

# Example mixed-precision applications

Mixed precision balances **performance** and **numerical accuracy**

- Deep Learning
  - Commonly uses 16 or 8-bit (even 1 bit!) floating-point numbers
- Fluid dynamics simulations
  - High precision only for critical parts, such as turbulence modeling
- Molecular Dynamics
  - Lower precision for long-range calculations
- Finite Element Analysis
  - Iterative methods to solve large linear systems

# Example radiation solver



# Floating-point format (IEEE 754)

- IEEE 754 standard is implemented in all architectures
- Floating-point number consists of three parts:
  - S: sign (+ or -)
  - M: mantissa/significand
  - E: exponent
- Floating-point number represented using exponential format:
  - $(-1)^S \times M \times 2^E$
  - Example:  $+1.42 \times 2^3$  means  $S=+1$ ,  $M=1.42$ ,  $E=3$
  - Where  $1 \leq M < 2$ , which makes representation unique
  - There are also non-normal numbers: NaN, Inf, subnormal


# Floating-point format (IEEE 754)

- Sign bit (1 bit)
- Mantissa/significand (A bits)
  - Determines number of **significant digits**
  - Results rounded to number of decimal places
  - Example: A=23 means ~7 decimal places
- Exponent (B bits)
  - Determines **range** of numbers
  - Numbers outside range become *zero* or *infinity*
  - Example: B=8 means range is  $\sim 10^{-38}$  to  $\sim 10^{38}$
- Total size: **1 + A + B bits**

Type	$\sqrt{2}$
A=52 (Float64)	1.41421356237309
A=23 (Float32)	1.414213
A=10 (Float16)	1.414
A=2 (Float8)	1.5




# FP types: double



- **double** precision (64 bits) prevalent in scientific computing
- GPUs typically slow on double arithmetic
  - Except the scientific/datacenter-rated GPUs

Type name	Total bits	Exponent bits	Significant bits	Smallest normal	Biggest normal	Decimal places	1+Epsilon
double	64	11	52	2.2e-308	1.8e+308	15	1 + 2.22e-16


# FP types: float



- **Single** precision (32 bits) balances accuracy and throughput
- Widely used in graphics and general GPU applications

Type name	Total bits	Exponent bits	Significant bits	Smallest normal	Biggest normal	Decimal places	1+Epsilon
float	32	8	23	1.2e-38	3.4e+38	6	1.000000119

# FP types: half




- Introduced with NVIDIA's Pascal architecture (2016)
- Double computational throughput of float
- Limited range, reasonable accuracy

Type name	Total bits	Exponent bits	Significant bits	Smallest normal	Biggest normal	Decimal places	1+Epsilon
half	16	5	10	0.000061	65536	3	1.00097



# FP types: bfloat




- "Brain" Floating-point. Introduced by Google Brain project
- Introduced with NVIDIA's Ampere architecture (2020)
- Large range, limited accuracy


Type name	Total bits	Exponent bits	Significant bits	Smallest normal	Biggest normal	Decimal places	1+Epsilon
bfloat	16	8	7	1.2e-38	3.4e+38	2	1.00781

# FP types: 8-bit floats

**fp8\_e5m2**



**fp8\_e4m3**



- Introduced with NVIDIA's Hopper architecture (2022)
- Two flavors: 5+2 bits or 4+3 bits
- No arithmetic functions, only conversions

Type name	Total bits	Exponent bits	Significant bits	Smallest normal	Biggest normal	Decimal places	1+Epsilon
fp8_e4m3	8	4	3	0.015625	256	1	1.125
fp8_e5m2	8	5	2	0.000061	65536	0	1.25

# Floating-point overview

Type name	Total bits	Exponent bits	Significant bits	Smallest normal	Biggest normal	Decimal places	1+Epsilon
double	64	11	52	2.2e-308	1.8e+308	15	1 + 2.22e-16
float	32	8	23	1.2e-38	3.4e+38	6	1.000000119
half	16	5	10	0.000061	65536	3	1.00097
bfloat	16	8	7	1.2e-38	3.4e+38	2	1.00781
fp8_e4m3	8	4	3	0.015625	256	1	1.125
fp8_e5m2	8	5	2	0.000061	65536	<1	1.25

# Mixed precision in practice for CUDA

- Create type aliases in kernels
  - C: use preprocess `#define`
  - C++: use template parameters
- Available data types in CUDA
  - `double` and `float` are predefined
  - `__half` found in `<cuda_fp16.h>`
  - `__nv_bfloat16` found in `<cuda_bf16.h>`
  - `__nv_fp8_eXmY` found in `<cuda_fp8.h>`

# Example kernel

```
__global__ void vector_add(  
    int n,  
    const float* A,  
    const float* B,  
    float* C  
) {  
    int i = threadIdx.x + blockIdx.x * blockDim.x;  
    if (i < n)  
        C[i] = A[i] + B[i];  
}
```

# Example kernel

```
#define TYPE_A float
#define TYPE_B float
#define TYPE_C float

__global__ void vector_add(
    int n,
    const TYPE_A* A,
    const TYPE_B* B,
    TYPE_C* C
) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < n)
        C[i] = A[i] + B[i];
}
```

# Example kernel

```
#include <cuda_fp16.h>

#define TYPE_A __half
#define TYPE_B float
#define TYPE_C float

__global__ void vector_add(
    int n,
    const TYPE_A* A,
    const TYPE_B* B,
    TYPE_C* C
) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < n)
        C[i] = A[i] + B[i];
}
```

## Example kernel

```
#include <cuda_fp16.h>

#define TYPE_A __half
#define TYPE_B float
#define TYPE_C float

__global__ void vector_add(
    int n,
    const TYPE_A* A,
    const TYPE_B* B,
    TYPE_C* C
) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < n)
        C[i] = A[i] + B[i];
}
```

**Does not compile!** 😞

kernel.cu(15): error: no operator "+" matches  
these operands  
operand types are: \_\_half + float



## Example kernel

```
#include <cuda_fp16.h>

#define TYPE_A __half
#define TYPE_B float
#define TYPE_C float

__global__ void vector_add(
    int n,
    const TYPE_A* A,
    const TYPE_B* B,
    TYPE_C* C
) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < n)
        C[i] = A[i] + B[i];
}
```

**Does not compile!** 😞

kernel.cu(15): error: no operator "+" matches  
these operands  
operand types are: \_\_half + float

# Mixed precision programming challenges

- No type promotion
  - Cannot mix types in binary operations
- Some operations require intrinsics
  - `__hdiv()`, `__hsin()`, `__hfmad()`
- Missing operations
  - No `__htan()`?
- Missing or awkward type conversion
  - `__nv_cvt_bfloat16raw2_to_fp8x2`
  - No fp8 to double?
  - No half to bfloat16?

```
__global__ void kernel(const __half* input, float constant, float* output) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    __half in0 = input[2 * i + 0];  
    __half in1 = input[2 * i + 1];  
    __half2 a = __halves2half2(in0, in1);  
    float b = float(constant);  
    __half c = __float2half(b);  
    __half2 d = __half2half2(c);  
    __half2 e = __hadd2(a, d);  
    __half f = __low2half(e);  
    __half g = __high2half(e);  
    float out0 = __half2float(f);  
    float out1 = __half2float(g);  
    output[2 * i + 0] = out0;  
    output[2 * i + 1] = out1;  
}
```

```
__global__ void kernel(const __half* input, float constant, float* output) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    __half in0 = input[2 * i + 0];
    __half in1 = input[2 * i + 1];
    __half2 a = __halves2half2(in0, in1);
    float b = float(constant);
    __half c = __float2half(b);
    __half2 d = __half2half2(c);
    __half2 e = __hadd2(a, d);
    __half f = __low2half(e);
    __half g = __high2half(e);
    float out0 = __half2float(f);
    float out1 = __half2float(g);
    output[2 * i + 0] = out0;
    output[2 * i + 1] = out1;
}
```



```
#include "kernel_float.h"
namespace kf = kernel_float;
```

```
__global__ void kernel(const kf::vec<half, 2>* input, float constant, kf::vec<float, 2>* output) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    output[i] = input[i] + kf::cast<half>(constant);
}
```




# Kernel Float

[https://github.com/KernelTuner/kernel\\_float](https://github.com/KernelTuner/kernel_float)

# Kernel Float

- Header-only C++ library to simplify mixed precision GPU programming
- Offers single type: `vec<T, N>`
  - N elements of type T
  - Auto selects optimal storage format
- Offers all mathematical operations
  - Auto selects best intrinsic
  - Fallback to single precision for missing operations



## Kernel Float

[Github](#)
[Docs](#)
[Build](#)
[Contribute](#)
[Blog](#)
[API](#)
[Sign Up](#)
[Sign In](#)
[Stars](#)

Kernel Float is a header-only library for CUDA that simplifies working with vector types and reduced precision floating-point arithmetic in GPU code.

### Summary

CUDA natively offers several reduced precision floating-point types (e.g., `_nv_half_t`, `_nv_bfloat16_t`, `_nv_f16_t`, `_nv_f8_t`) and vector types (e.g., `_nv_half2_t`, `_nv_half4_t`, `_nv_float2_t`, `_nv_float4_t`). However, working with these types is cumbersome: mathematical operations require intrinsics (e.g., `_nv_half2_add` performs addition for `_nv_half2_t`), type conversion is awkward (e.g., `_nv_cvt_half2_to_float2` converts `float2` to `float2`), and some functionality is missing (e.g., one cannot convert a `_nv_half_t` to `_nv_bfloat16_t`).

Kernel Float resolves this by offering a single data type `kernel_float::vec<T, N>` that stores `N` elements of type `T`. Internally, the data is stored as a fixed-sized array of elements. Operator overloading (like `+`, `*`, `^`, `&&`) has been implemented such that the most optimal intrinsic for the available types is selected automatically. Many mathematical functions (like `log`, `exp`, `sin`) and common operations (such as `sum`, `range`, `for_each`) are also available.

By using this library, developers can avoid the complexity of working with reduced precision floating-point types in CUDA and focus on their applications.

### Features

In a nutshell, Kernel Float offers the following features:

- Single type `vec<T, N>` that unifies all vector types.
- Operator overloading to simplify programming.
- Support for half (16 bit) floating-point arithmetic, with a fallback to single precision for unsupported operations.
- Support for quarter (8 bit) floating-point types.
- Easy integration as a single header file.
- Written for C++17.
- Compatible with NVCC (NVIDIA Compiler) and NVRTC (NVIDIA Runtime Compilation).

### Example

Check out the `examples` directory for some examples.

Below shows a simple example of a CUDA kernel that adds a `constant` to the `input` array and writes the results to the `output` array. Each thread processes two elements. Notice how easy it would be to change the precision (for example, `double` to `half`) or the vector size (for example, 4 instead of 2 items per thread).

```

#include "kernel_float.h"
namespace kf = kernel_float;

__global__ void kernel(const kf::vec_half, 2* input, float constant, kf::vec_float, 2* output

```

# Example kernel

```
#define TYPE_A float
#define TYPE_B float
#define TYPE_C float

__global__ void vector_add(
    int n,
    const TYPE_A* A,
    const TYPE_B* B,
    TYPE_C* C
) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < n)
        C[i] = A[i] + B[i];
}
```

# Example kernel

```
#include "kernel_float.h"

#define TYPE_A float
#define TYPE_B float
#define TYPE_C float

__global__ void vector_add(
    int n,
    const kernel_float::vec<TYPE_A, 1>* A,
    const kernel_float::vec<TYPE_B, 1>* B,
    kernel_float::vec<TYPE_C, 1>* C
) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < n)
        C[i] = A[i] + B[i];
}
```

# Example kernel

```
#include "kernel_float.h"

#define TYPE_A __half
#define TYPE_B float
#define TYPE_C float

__global__ void vector_add(
    int n,
    const kernel_float::vec<TYPE_A, 1>* A,
    const kernel_float::vec<TYPE_B, 1>* B,
    kernel_float::vec<TYPE_C, 1>* C
) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < n)
        C[i] = A[i] + B[i];
}
```



# Vectorization

- Kernel Float automatically uses **vector** intrinsics
  - Requires using `kernel_float::vec<T, N>` with  $N \geq 2$
- Several types benefit from vectorization!
  - `half` and `bfloat` require vectorized intrinsics for high throughput
  - Vectorized memory operations
  - Vectorized integer operations
  - ...

# Example kernel

```
#include "kernel_float.h"

#define TYPE_A float
#define TYPE_B float
#define TYPE_C __half
#define VECTOR_SIZE 1

__global__ void vector_add(
    int n,
    const kernel_float::vec<TYPE_A, VECTOR_SIZE>* A,
    const kernel_float::vec<TYPE_B, VECTOR_SIZE>* B,
    kernel_float::vec<TYPE_C, VECTOR_SIZE>* C
) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i * VECTOR_SIZE < n)
        C[i] = A[i] + B[i];
}
```

# Example kernel

```
#include "kernel_float.h"

#define TYPE_A float
#define TYPE_B float
#define TYPE_C half
#define VECTOR_SIZE 2

__global__ void vector_add(
    int n,
    const kernel_float::vec<TYPE_A, VECTOR_SIZE>* A,
    const kernel_float::vec<TYPE_B, VECTOR_SIZE>* B,
    kernel_float::vec<TYPE_C, VECTOR_SIZE>* C
) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i * VECTOR_SIZE < n)
        C[i] = A[i] + B[i];
}
```

# Tuning Problem

- **Accuracy vs performance** trade-off
  - What type should we use for each variable?
  - Ideally want **high performance** with **low error**
- Variables datatypes and **kernel parameters** both affect performance
  - Usually heavily **intertwined**, we cannot tune them separately
- Leads to large **search-space**, for example:
  - 10 variables and 4 precision levels:  $4^{10} = 1$  million options
  - 8 parameters with each 6 options:  $6^8 = 1$  million options
  - Total: 1 trillion configurations!

# Kernel Tuner

Kernel Tuner offers native support for *accuracy tuning*

- **Step 1:** Add tunable floating-point types as **tuning parameters**
- **Step 2:** Wrap inputs/outputs in **TunablePrecision** objects
- **Step 3:** provide **reference output** as answer
- **Step 4:** Add **AccuracyObserver**

See the example:

- [examples/cuda/accuracy.py](#)

# Kernel Tuner Example

```
size = 100000000
n = numpy.int32(size)
a = numpy.random.randn(size)
b = numpy.random.randn(size)
c = numpy.zeros_like(b)

args = [n,
        TunablePrecision("float_type", a),
        TunablePrecision("float_type", b),
        TunablePrecision("float_type", c)]

answer = [None, None, None, a + b]

tune_params = dict()
tune_params["block_size_x"] = [32, 64, 128, 256, 512, 1024]
tune_params["float_type"] = ["float", "double", "half"]

observers = [AccuracyObserver("RMSE")]

results, env = tune_kernel("vector_add", kernel_string,
                           size, args, tune_params, answer=answer,
                           observers=observers, lang="CUDA")
```

# Kernel Tuner Example

```
size = 100000000
n = numpy.int32(size)
a = numpy.random.randn(size)
b = numpy.random.randn(size)
c = numpy.zeros_like(b)

args = [n,
        TunablePrecision("float_type", a),
        TunablePrecision("float_type", b),
        TunablePrecision("float_type", c)]

answer = [None, None, None, a + b]

tune_params = dict()
tune_params["block_size_x"] = [32, 64, 128, 256, 512, 1024]
tune_params["float_type"] = ["float", "double", "half"]

observers = [AccuracyObserver("RMSE")]

results, env = tune_kernel("vector_add", kernel_string,
                           size, args, tune_params, answer=answer,
                           observers=observers, lang="CUDA")
```

# Kernel Tuner Example

```
size = 100000000
n = numpy.int32(size)
a = numpy.random.randn(size)
b = numpy.random.randn(size)
c = numpy.zeros_like(b)

args = [n,
        TunablePrecision("float_type", a),
        TunablePrecision("float_type", b),
        TunablePrecision("float_type", c)]

answer = [None, None, None, a + b]

tune_params = dict()
tune_params["block_size_x"] = [32, 64, 128, 256, 512, 1024]
tune_params["float_type"] = ["float", "double", "half"]

observers = [AccuracyObserver("RMSE")]

results, env = tune_kernel("vector_add", kernel_string,
                           size, args, tune_params, answer=answer,
                           observers=observers, lang="CUDA")
```



## Tunable and TunablePrecision

- The **TunablePrecision** wrapper tells Kernel Tuner that type of input/output arguments depends on a **tunable parameter**
- **Before** benchmarking, data **converted** to provided data types
- **During** benchmarking, kernel is passed **pointer** of correct data type
- [Advanced] The general **Tunable** object allows arbitrary conversions

# Kernel Tuner Example

```
size = 100000000
n = numpy.int32(size)
a = numpy.random.randn(size)
b = numpy.random.randn(size)
c = numpy.zeros_like(b)

args = [n,
        TunablePrecision("float_type", a),
        TunablePrecision("float_type", b),
        TunablePrecision("float_type", c)]

answer = [None, None, None, a + b]

tune_params = dict()
tune_params["block_size_x"] = [32, 64, 128, 256, 512, 1024]
tune_params["float_type"] = ["float", "double", "half"]

observers = [AccuracyObserver("RMSE")]

results, env = tune_kernel("tuner_01", kernel_string,
                           size, args, tune_params, answer=answer,
                           observers=observers, lang="CUDA")
```

# Kernel Tuner Example

```
size = 100000000
n = numpy.int32(size)
a = numpy.random.randn(size)
b = numpy.random.randn(size)
c = numpy.zeros_like(b)

args = [n,
        TunablePrecision("float_type", a),
        TunablePrecision("float_type", b),
        TunablePrecision("float_type", c)]

answer = [None, None, None, a + b]

tune_params = dict()
tune_params["block_size_x"] = [32, 64, 128, 256, 512, 1024]
tune_params["float_type"] = ["float", "double", "half"]

observers = [AccuracyObserver("RMSE")]

results, env = tune_kernel("vector_add", kernel_string,
                           size, args, tune_params, answer=answer,
                           observers=observers, lang="CUDA")
```

# Error metrics

- The **AccuracyObserver** measures the error and adds a metric
- Supports 10+ well-known metrics
  - Root mean square error (RMSE)
  - Mean relative error (rel)
  - Mean absolute error (abs)
  - Maximum relative error (max)
  - ....
  - Custom metrics are also possible!
- The best error metric is **application-dependent**
- Compatible with other observers!


# Hands-on

# Second hands-on

- The second hands-on notebook is:
  - [https://github.com/KernelTuner/kernel\\_tuner\\_tutorial/blob/master/hands-on/esiwace3/05\\_Mixed\\_precision\\_programming.ipynb](https://github.com/KernelTuner/kernel_tuner_tutorial/blob/master/hands-on/esiwace3/05_Mixed_precision_programming.ipynb)

- The goal of this hands-on is to:
  - Tune a signal convolution kernel with **mixed precision types**
  - Experiment with the **accuracy-performance** trade-off

- Please follow the instructions in the Notebook
- Feel free to ask questions to instructors and mentors



```
# The tunable types. Currently, the code is only tuned for double
tune_params = dict()
tune_params["OUTPUT_TYPE"] = ["double"]
tune_params["INPUT_TYPE"] = ["double"]
tune_params["FILTER_TYPE"] = ["double"]
tune_params["ACCUM_TYPE"] = ["double"]

# Other tunable parameters
tune_params["block_size_x"] = [128, 256]
tune_params["VECTOR_SIZE"] = [1, 2, 4]
tune_params["PREFETCH_INPUT"] = [0, 1]
tune_params["UNROLL_LOOP"] = [0, 1]
```

# **Optimizing GPU Core Clock Frequency**

## Measuring power consumption with NVML

NVML (the Nvidia Management Library) can observe GPU temperature, core and memory clocks, core voltage, and power

### Advantages:


- Highly available

### Disadvantages:

- Returns time-averaged power, not instantaneous power consumption
- Limited time resolution

### Current solution:

- Measure power while continuously running the kernel for one second






# PowerSensor2

## Pros:

- Instantaneous power readings
- Time resolution: 2.8 KHz
- Open source: <https://gitlab.com/astron-misc/PowerSensor>

## Cons:

- Some assembly required
  - You need to build the hardware!



Supported in Kernel Tuner, using [PowerSensorObserver](#)

# NVMLObserver - Nvidia Management Library (NVML)

Allows to measure several quantities during tuning:

- Power consumption, core frequency, core voltage, memory frequency, GPU temperature, and energy consumption

Provides an interface within Kernel Tuner to NVML:

- Enables new tunable parameters:
  - `nvml_pwr_limit`: try out different power limits
  - `nvml_gr_clock`: set the GPU core clock frequency
  - `nvml_mem_clock`: set the GPU memory clock frequency
  - **Setting these requires root privileges**

## NVMLObserver example

```
...  
tune_params["nvm1_pwr_limit"] = [250, 225, 200, 175]  
  
nvmlobserver = NVMLObserver(["nvm1_energy", "temperature"])  
  
metrics = OrderedDict()  
metrics["GFLOPS/W"] = lambda p: (size/1e9) / p["nvm1_energy"]  
  
results, env = tune_kernel("vector_add", kernel_string, size, args,  
                           tune_params, observers=[nvmlobserver],  
                           metrics=metrics, iterations=32)
```

# Setup GPU power limits and core and memory clock frequencies

- Kernel Tuner has helper functions to setup tunable parameters:

In `kernel_tuner.observers.nvml`:

- `get_nvml_pwr_limits(device, n=None, quiet=False)`:
  - Device is the device ordinal as reported by `nvidia-smi`
  - `n` is the number of evenly-spaced values to tune
    - if unspecified returns values spaced 5 Watts apart
- `get_nvml_gr_clocks(device, n=None, quiet=False)`:
  - `n` is the number of evenly-spaced values to tune
    - If unspecified, all supported core clocks are returned

# Setup core clock frequencies for tuning

```
tune_params["block_size_x"] = [128, 256, 512, 1024]
...

tune_params.update(get_nvml_gr_clocks(0, 7)) # get 7 core frequencies supported on device 0

nvmlobserver = NVMLObserver(["nvml_energy", "temperature"])

metrics = OrderedDict()
metrics["GFLOPS/W"] = lambda p: (size/1e9) / p["nvml_energy"]

results, env = tune_kernel("vector_add", kernel_string, size, args,
                           tune_params, observers=[nvmlobserver],
                           metrics=metrics, iterations=32)
```

# Power limit vs frequency tuning

Many tunable parameters affect compute performance and/or energy efficiency


But we can also:

- Limit the GPU clock frequency, allow GPU to vary power consumption
- Limit the GPU power consumption, allow GPU to determine clock frequency

Both methods unfortunately require root privileges for the latest generations of Nvidia GPUs

# Frequency-power relation

Tuning CLBlast GEMM using frequency or power limit tuning



# Frequency tuning vs power capping

## Advantages of power capping:

- Potentially more effective, GPU may also lower memory clock
- Reliable method in face of limited power

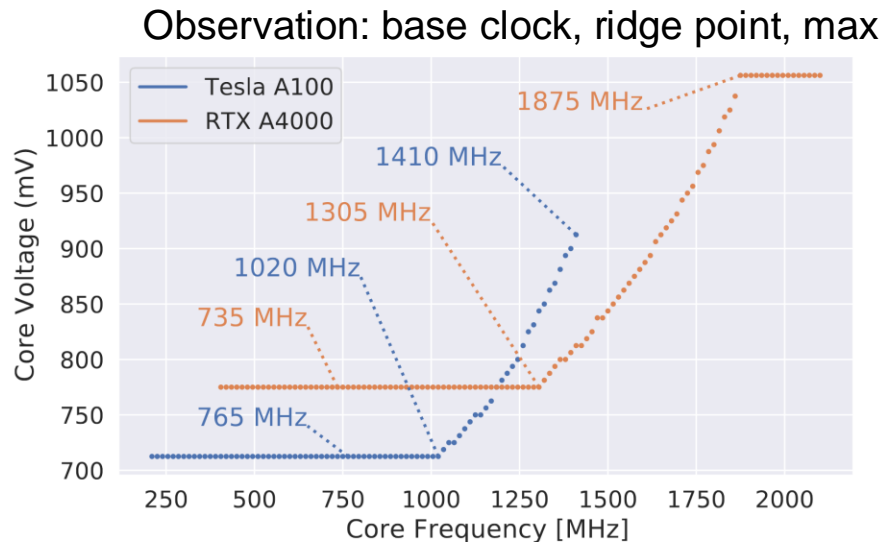
## Advantages of frequency tuning:

- Especially on A100, frequency tuning enables a wider power range
- Fixing the clock frequency also improves measurement stability




# Frequency Voltage relation

- GPUs rapidly ramp up voltage when clock frequency increases beyond a certain point
- This point appears to be a sweet spot in the trade-off between energy consumption and compute performance
- We call this point the 'ridge point'




# A simple power consumption model

- Not every GPU reports core voltages, but we can estimate the voltage using a simple power model
- When we fix all parameters and vary the clock frequency, we can approximate power consumption using:
$$P_{load} = \min(P_{max}, P_{idle} + \alpha * f * v^2)$$
- And identify the GPUs 'ridge point' frequency in this way



# Model-steered auto-tuning

- Use performance model to limit the frequency range for tuning
- Reduces the search space by ~80% on average



## Custom tuning objectives

- By default, Kernel Tuner's optimization strategy minimizes time
- But there is also support for using a custom tuning **objective**
- The objective can be any observed quantity or user-defined metric

```
metrics["GFLOPS/W"] = lambda p: (size/1e9) / p["nvm1_energy"]
```

```
results, env = tune_kernel("vector_add", kernel_string, size, args,  
                           tune_params, observers=[nvmlobserver],  
                           metrics=metrics, iterations=32,  
                           objective="GFLOPS/W")
```

# Hands-on

## Third hands-on

- The third hands-on notebook is:
  - [https://github.com/KernelTuner/kernel\\_tuner\\_tutorial/blob/master/hands-on/esiwace3/06\\_Energy\\_Efficient\\_Computing.ipynb](https://github.com/KernelTuner/kernel_tuner_tutorial/blob/master/hands-on/esiwace3/06_Energy_Efficient_Computing.ipynb)
- The goal of this hands-on is to:
  - Tune a kernel to minimize the execution time or the **energy consumption**
  - Use an optimization strategy
  - Compare different **energy optimization strategies**
- Please follow the instructions in the Notebook
- Feel free to ask questions to instructors and mentors

# Closing Remarks

# How do we create energy efficient GPU applications?

Three strategies for energy efficient GPU Computing:

1. Use for shorter amount of time → Optimize application performance
2. Minimize data movements → Lower/mixed precision techniques
3. Optimize device settings → Optimize clock frequency



# Learning objectives

- Understand the energy footprint of computing
- Optimize applications for performance to reduce energy consumption
- Reduce data movement with mixed-precision techniques
- Tune GPU core frequencies to find the most energy-efficient configuration

# Contributions are welcome!

- Contributions can come in many forms: tweets, blog posts, issues, pull requests
- Before making larger changes, please create an issue to discuss
- For the full contribution guide, please see:  
[https://kerneltuner.github.io/kernel\\_tuner/stable/contributing.html](https://kerneltuner.github.io/kernel_tuner/stable/contributing.html)

# Related publications

- **Kernel Launcher: C++ library for creating optimal-performance portable CUDA applications**  
S. Heldens, B. van Werkhoven  
*International Workshop on Automatic Performance Tuning (iWAPT2023) co-located with IPDPS 2023*
- **Optimization Techniques for GPU Programming**  
Pieter Hijma, Stijn Heldens, Alessio Sclocco, Ben van Werkhoven, and Henri Bal  
*ACM Computing surveys 2023*
- **Going green: optimizing GPUs for energy efficiency through model-steered auto-tuning**  
Richard Schoonhoven, Bram Veenboer, Ben van Werkhoven, K. Joost Batenburg  
*International Workshop on Performance Modeling, Benchmarking and Simulation of High-Performance Computer Systems (PMBS) at Supercomputing (SC22) 2022*
- **Bayesian Optimization for auto-tuning GPU kernels**  
F.J. Willemsen, R.V. van Nieuwpoort, B. van Werkhoven  
*International Workshop on Performance Modeling, Benchmarking and Simulation of High Performance Computer Systems (PMBS) at Supercomputing (SC21) 2021*
- **Kernel Tuner: A search-optimizing GPU code auto-tuner**  
B. van Werkhoven  
*Future Generation Computer Systems 2019*



## Interested in getting in touch?



Website: [www.esiwace.eu](http://www.esiwace.eu)



Twitter: <https://twitter.com/esiwace>



YouTube: <https://www.youtube.com/@esiwace880>

ESiWACE is on Zenodo, the Open Access repository for scientific results  
<https://zenodo.org/communities/esiwace>



**EuroHPC**  
Joint Undertaking

*Funded by the European Union. This work has received funding from the European High Performance Computing Joint Undertaking (JU) under grant agreement No 101093054.*



Co-funded by  
the European Union

netherlands  
**eScience** center